

# COMP3702 Tutorial 8

Matt Choy | [matthew.choy@uq.edu.au](mailto:matthew.choy@uq.edu.au)

# COMP3702 Tutorial 8

Matt Choy | [matthew.choy@uq.edu.au](mailto:matthew.choy@uq.edu.au)

<https://github.com/mattpchoy/comp3702-tutorials>

# Monte Carlo Tree Search

- MCTS belongs to a family of algorithms for fast planning in online settings
- The longer MCTS is allowed to run for, the higher quality the resulting policy is (on average)
- Rather than iterating over every connection in the state graph representation, we build a tree (subset of state graph), and prioritise the important states.
- The MCTS algorithm, consists of four components:
  1. Selection
  2. Expansion
  3. Simulation
  4. Back-Propagation
- At each node, we store the statistics:
  - $Q(s, a)$  for each action in  $A$  – this is the average reward for  $(s, a)$  over all of our trials
  - $N(s, a)$  for each action in  $A$  – this is the number of times action  $a$  has been performed from state  $s$ .
  - $N(s)$  this is the number of the times this state has been visited with any action performed.

# Monte Carlo Tree Search - Selection

- Given that we are at a current state, how do we choose what action to perform?
- Aim to compromise between:
  - Exploration (visiting under-explored branches) and
  - Exploitation (Visiting branches with higher average reward)

# Monte Carlo Tree Search - Selection

- Given that we are at a current state, how do we choose what action to perform?
- Aim to compromise between:
  - Exploration (visiting under-explored branches) and
  - Exploitation (Visiting branches with higher average reward)
- Selection strategies:
  - Random Choice – if any actions have never been tried, choose an untried action at uniform random (tends to have better performance than trying to choose actions in a fixed order).
  - Epsilon-Greedy – Choose the highest Q-action with probability  $\epsilon$  and all other actions equally likely (with  $\frac{1-\epsilon}{n}$  probability, where  $n$  is the number of actions that can be performed).
  - UCB – Compute a confidence interval for the true average reward, based on the number of trials and choose the action with highest UCB

# Monte Carlo Tree Search - Expansion

- Convert a leaf node into a non-leaf node
- When a leaf node is reached:
  - Set  $N(s) \leftarrow 1$  (initialise the node count to 1)
  - Estimate  $V$  (the future expected value of the state) via simulation

# Monte Carlo Tree Search - Simulation

- Estimate the future expected value of a state without building up a tree
- Random Roll-out Choose actions at random until some maximum horizon is reached, keeping a running total of the reward
  - E.g., look 20 steps in the future, discounted by  $\gamma$
  - This is not necessarily close to the optimum value for the state, but it is a “good enough” estimate
  - It indicates whether the state leads to potentially dangerous states, or states with high reward.
- Can average this over a number of random rollouts.
- Can use a heuristic to choose actions during roll-out rather than choosing purely at random
- Return the estimated value  $V$

# Monte Carlo Tree Search - Backpropagation

- We want to use the results from our simulations and update our node statistics and values.
- Update the average total discounted reward and node/action counts for each branch visited.
- Let  $s_t$  be the state visited at time step  $t$ , and let  $a_t$  be the action performed at time step  $t$
- At time step  $t$ , the total discounted future reward is given by the following equation:

$$R_t = r(s_t) + \gamma^1 r(s_{t+1}) + \dots + \gamma^n r(s_{t+n}) + \gamma^{n+1} V$$

- For all time steps, we compute the Q-value, and update node statistics.

$$Q(s_t, a_t) \leftarrow \frac{N(s_t)Q(s_t, a_t) + R_t}{N(s_t) + 1}$$

$$N(s_t, a_t) \leftarrow N(s_t, a_t) + 1$$

$$N(s_t) \leftarrow N(s_t) + 1$$



# MCTS Implementation – Iteratively

- Create a Tree Node class that stores:
  - $N(s)$
  - $Q(s, a)$  and  $N(s, a)$  for each available action
  - Stores a list of child nodes for each available actions
  - Stores a reference to the parent node
- `mcts_search(current_state)`
  - `node ← current_state`
  - While the node is not a leaf node, select an action and sample a next state (and set `node ← next_state`)
  - Expand the leaf node and estimate the value via simulation
    - Create a new `TreeNode` instance for it
  - While the node doesn't have a parent, update  $Q(s,a)$ ,  $N(s, a)$  and  $N(a)$ 
    - And set `Node ← node.parent`
    - This is our backpropagation step (where we move backward in time, and update values)
    - We keep repeating this step until we reach the root node.

# MCTS Implementation – Recursively

- Dictionaries are used to store node statistics

$N_s[s]$  → Number of times a state “s” has been visited

$N_{s,a}[(s, a)]$  → Number of times an action “a” has been performed from state “s”

$Q_{s,a}[(s, a)]$  → Average reward from performing action a at state s

- `mcts_search(current_state)`
  - If the current state is a leaf node, estimate the value V from simulation and return the value (recursion base case)
  - Otherwise, select an action for the current state using our dictionaries
    - Action can be selected using epsilon-greedy or UCB
  - Sample the outcome of the next state, and set
$$V = \text{immediate\_reward} + \gamma \times \text{mcts\_search}(\text{next\_state})$$
  - Increment  $N_s[s]$  and  $N_{s,a}[(s, a)]$  and update  $Q_{s,a}[(s, a)]$  using the value V
  - Return the value V (so that the next level above can use the value)

# MCTS Implementation

- For both approaches, the `mcts_select_action` method should:
  - Call `mcts_search(current_state)` while time / memory / iteration limits are not reached
  - Action  $\leftarrow \operatorname{argmax}(Q(s,a))$  over all actions
  - Return the action